

Lecture 27: November 29

*Lecturer: Vijay Garg**Scribe: Chunheng Luo*

Question 1

Which of the following statements about blocks and threads in CUDA are **FALSE**?

- A. The maximum number of threads per block is determined by the hardware architecture of the GPU.
- B. The method `__syncthreads()` acts as a barrier to synchronize all threads within a block.
- C. Two threads located in different blocks can synchronize with each other within the kernel.
- D. Branches (e.g. if/else) in the kernel can potentially reduce parallelism between threads.

Solution: C

Question 2

Which of the following statements about shared and global memory in CUDA are **FALSE**?

- A. Global memory has the scope of the entire application.
- B. Shared memory is usually on-chip and faster than global memory.
- C. Global memory has to be declared outside the kernel.
- D. Shared memory cannot be accessed by the host.

Solution: C

Question 3

Which of the following statements about kernel functions in CUDA are **TRUE**?

- A. Kernel functions are called by the host and run on the device.
- B. The return type of kernel functions can only be `void`
- C. We can have multiple kernel invocations in an application.
- D. All of the above.

Solution: D

Question 4

In the work-optimal algorithm for computing array maximum discussed in class,

- A. the array of size n is partitioned into $(n/\log \log n)$ groups of size $O(\log \log n)$.
- B. sequential algorithm is run on each of the groups to get the maximum value of each group.
- C. the doubly logarithmic height tree algorithm is used to get the maximum value of the group maximums.
- D. All of the above.

Solution: D

Question 5

In the work-optimal algorithm for merging two sorted arrays discussed in class,

- A. the ranks of splitters are computed using the parallel rank-computation algorithm.
- B. sublists are merged using the sequential merging algorithm.
- C. the overall time complexity is $O(\log n)$.
- D. All of the above.

Solution: D

Question 6

In the list ranking algorithm discussed in class, what technique is used to reduce the length of the longest path in the linked structure in each iteration?

Solution: Pointer jumping

Question 7

In the breaking symmetry algorithm discussed in class, what is a possible method to guarantee that initially all nodes have valid coloring?

Solution: Each node could initially have different color by using its unique identifier.

Question 8

Briefly describe the upward sweep and downward sweep procedures in the Blelloch Scan algorithm.

Solution:

Upward: $sum[v] = sum[Left[v]] + sum[Right[v]]$;

Downward: $scan[Left[v]] = scan[v]$; $scan[Right[v]] = sum[Left[v]] + scan[v]$;